

Inference Latency on adult (lower=better)
GBDTs + MLP-family DL are instant; TabPFN & STab are ~1000x+ slower to serve

